

Experiment No.5

- **Title:-** Implement Bag-of-Words (Count Occurrence, Normalized Count Occurrence), TF-IDF, and Word2Vec Embeddings.
- **Aim:-** To Implement Bag-of-Words (Count Occurrence, Normalized Count Occurrence), TF-IDF, and Word2Vec Embeddings.

- **Objective:-**

1. Understand the concept of Bag-of-Words (BoW) and its limitations.
2. Implement Bag-of-Words (count occurrence and normalized form) using Python.
3. Apply TF-IDF to assign importance weights to words.
4. Generate dense vector representations using Word2Vec embeddings.
5. Compare sparse vs. dense representations of text data.

- **Procedure and Flowchart:-**

Step 1: Install required Python libraries:

```
pip install scikit-learn gensim
```

Step 2: Import necessary modules from sklearn and gensim.

Step 3: Define a sample text corpus.

Step 4: Apply Bag-of-Words (Count Occurrence).

Step 5: Apply Normalized Bag-of-Words (Term Frequency).

Step 6: Apply TF-IDF transformation.

Step 7: Train Word2Vec model on the text data.

Step 8: Display results for each method and analyze differences.

Name :- Prathamesh Arvind Jadhav

Roll No:- 4059

- **Python code:-**

```
import numpy as np

import pandas as pd

from sklearn.feature_extraction.text import CountVectorizer, TfidfVectorizer

from gensim.models import Word2Vec

from nltk.tokenize import word_tokenize

import matplotlib.pyplot as plt

import seaborn as sns

# Sample data (replace with your dataset)

corpus = [

    "Natural language processing enables computers to understand human language",

    "Word embeddings like Word2Vec represent words in vector space",

    "Bag of words and TF-IDF are traditional NLP approaches",

    "Deep learning models improve performance in NLP tasks"

]

# 1. Bag-of-Words (Count Occurrence)

count_vectorizer = CountVectorizer()

count_matrix = count_vectorizer.fit_transform(corpus)

print("Vocabulary:", count_vectorizer.get_feature_names_out())

print("\nBag-of-Words (Count Occurrence):")

print(pd.DataFrame(count_matrix.toarray(),

columns=count_vectorizer.get_feature_names_out()))

# Visualization: Bag-of-Words

plt.figure(figsize=(8, 4))

sns.heatmap(pd.DataFrame(count_matrix.toarray(),

columns=count_vectorizer.get_feature_names_out()), annot=True, cmap="Blues")

plt.title("Bag-of-Words (Count Occurrence)")
```

Name :- Prathamesh Arvind Jadhav

Roll No:- 4059

```
plt.xlabel("Words")
```

```
plt.ylabel("Documents")
```

```
plt.tight_layout()
```

```
plt.show()
```

```
# 2. Normalized Count Occurrence (Frequency)
```

```
normalized_count = count_matrix.toarray().astype(float)
```

```
normalized_count = normalized_count / normalized_count.sum(axis=1, keepdims=True)
```

```
print("\nNormalized Count Occurrence (Term Frequency):")
```

```
print(pd.DataFrame(normalized_count, columns=count_vectorizer.get_feature_names_out()))
```

```
# Visualization: Normalized Count Occurrence
```

```
plt.figure(figsize=(8, 4))
```

```
sns.heatmap(pd.DataFrame(normalized_count,  
columns=count_vectorizer.get_feature_names_out()), annot=True, cmap="Greens")
```

```
plt.title("Normalized Count Occurrence (Term Frequency)")
```

```
plt.xlabel("Words")
```

```
plt.ylabel("Documents")
```

```
plt.tight_layout()
```

```
plt.show()
```

```
# 3. TF-IDF
```

```
tfidf_vectorizer = TfidfVectorizer()
```

```
tfidf_matrix = tfidf_vectorizer.fit_transform(corpus)
```

```
print("\nTF-IDF Matrix:")
```

```
print(pd.DataFrame(tfidf_matrix.toarray(), columns=tfidf_vectorizer.get_feature_names_out()))
```

```
# Visualization: TF-IDF
```

```
plt.figure(figsize=(8, 4))
```

```
sns.heatmap(pd.DataFrame(tfidf_matrix.toarray(),  
columns=tfidf_vectorizer.get_feature_names_out()), annot=True, cmap="Purples")
```

Name :- Prathamesh Arvind Jadhav

Roll No:- 4059

```
plt.title("TF-IDF Matrix")
```

```
plt.xlabel("Words")
```

```
plt.ylabel("Documents")
```

```
plt.tight_layout()
```

```
plt.show()
```

```
# 4. Word2Vec Embeddings
```

```
# Tokenize corpus
```

```
tokenized_corpus = [word_tokenize(doc.lower()) for doc in corpus]
```

```
# Train Word2Vec
```

```
w2v_model = Word2Vec(sentences=tokenized_corpus, vector_size=50, window=5,  
min_count=1, workers=4)
```

```
# Example: Get embedding of a word
```

```
word = "language"
```

```
print(f"\nWord2Vec embedding for '{word}':\n", w2v_model.wv[word])
```

```
# Example: Find most similar words
```

```
print("\nMost similar words to 'language':", w2v_model.wv.most_similar("language"))
```

- **Output:-**

Name :- Prathamesh Arvind Jadhav

Roll No:- 4059

```

PS D:\NLP\Exp5> python Exp5.py
Vocabulary: ['and', 'approaches', 'are', 'bag', 'computers', 'deep', 'embeddings', 'enables',
'human', 'idf', 'improve', 'in', 'language', 'learning', 'like', 'models',
'natural', 'nlp', 'of', 'performance', 'processing', 'represent', 'space',
'tasks', 'tf', 'to', 'traditional', 'understand', 'vector', 'word', 'word2vec',
'words']

Bag-of-Words (Count Occurrence):
  and approaches are bag computers deep ... traditional understand vector word word2vec
words
0 0 0 0 0 1 0 ... 0 1 0 0 0
0 0
1 0 0 0 0 0 0 ... 0 0 1 1 1
1 1
2 1 1 1 0 0 ... 1 0 0 0 0
1
3 0 0 0 0 1 ... 0 0 0 0 0
0

[4 rows x 32 columns]

Normalized Count Occurrence (Term Frequency):
  and approaches are bag computers deep ... traditional understand vector word w
ord2vec words
0 0.0 0.0 0.0 0.111111 0.000 ... 0.0 0.111111 0.000000 0.000000 0
.000000 0.000000
1 0.0 0.0 0.0 0.0 0.000000 0.000 ... 0.0 0.000000 0.111111 0.111111 0
.111111 0.111111
2 0.1 0.1 0.1 0.1 0.000000 0.000 ... 0.1 0.000000 0.000000 0.000000 0
.000000 0.100000
3 0.0 0.0 0.0 0.0 0.000000 0.125 ... 0.0 0.000000 0.000000 0.000000 0
.000000 0.000000

[4 rows x 32 columns]

```

```

TF-IDF Matrix:
  and approaches are bag computers ... understand vector word word2v
ec words
0 0.000000 0.000000 0.000000 0.000000 0.301511 ... 0.301511 0.000000 0.000000 0.0000
00 0.000000
1 0.000000 0.000000 0.000000 0.000000 0.000000 ... 0.000000 0.348299 0.348299 0.3482
99 0.274603
2 0.328919 0.328919 0.328919 0.328919 0.000000 ... 0.000000 0.000000 0.000000 0.0000
00 0.259324
3 0.000000 0.000000 0.000000 0.000000 0.000000 ... 0.000000 0.000000 0.000000 0.0000
00 0.000000

[4 rows x 32 columns]

Word2Vec embedding for 'language':
[-0.01631583 0.0089916 -0.00827415 0.00164907 0.01699724 -0.00892435
0.009035 -0.01357392 -0.00709698 0.01879702 -0.00315531 0.00064274
-0.00828126 -0.01536538 -0.00301602 0.00493959 -0.00177605 0.01106732
-0.00548595 0.00452013 0.01091159 0.01669191 -0.00290748 -0.01841629
0.0087411 0.00114357 0.01488382 -0.00162657 -0.00527683 -0.01750602
-0.00171311 0.00565313 0.01080286 0.01410531 -0.01140624 0.00371764
0.01217773 -0.0095961 -0.00621452 0.01359526 0.00326295 0.00037983
0.00694727 0.00043555 0.01923765 0.01012121 -0.01783478 -0.01408312
0.00180291 0.01278507]


Most similar words to 'language': [('natural', 0.37571486830711365), ('vector', 0.22977666556835175),
('models', 0.22082354128360748), ('approaches', 0.21908226609230042), ('bag', 0.16089819371700287),
('space', 0.14872416853904724), ('tasks', 0.12489066272974014), ('embeddings', 0.08060287684202194),
('enables', 0.07399576157331467), ('and', 0.05545045807957649)]

```

Name :- Prathamesh Arvind Jadhav

Roll No:- 4059

Figure 1



• Question and Answer:-

Q1. What is the difference between Bag-of-Words and TF-IDF?

Ans->

Aspect	Bag-of-Words (BoW)	TF-IDF
Weighting	Pure word counts/frequencies	Weighted by importance (TF × IDF)
Common words effect	High frequency → high importance (even for stopwords like <i>the</i>)	Penalized (low IDF)
Captures importance?	No	Yes
Example use cases	Simple text classification, spam detection	Information retrieval, search engines, NLP tasks needing meaningful features

Q2. Why is normalization important in Bag-of-Words representation?**Ans->**

Normalization in the Bag-of-Words (BoW) representation is important because it removes the bias caused by differences in document length. In BoW, documents are represented using raw word counts, which means longer documents naturally produce higher values even if the relative importance of words is the same as in shorter documents. Without normalization, this leads to unfair comparisons, where algorithms may mistakenly give more weight to longer texts. By normalizing, we scale the word counts so that documents can be compared on the basis of word distribution rather than size. For example, L1 normalization converts counts into relative frequencies, while L2 normalization ensures the vector length is 1, making cosine similarity and other distance-based measures more meaningful. In short, normalization makes the BoW features consistent, comparable, and more useful for machine learning models.

Q3. Explain how Word2Vec captures semantic meaning compared to BoW?**Ans->****Bag-of-Words (BoW)**

- **How it works:** Represents text as a sparse vector of word counts or frequencies.
- **Problem:**
 - Words are treated as independent; no relationship between them.
 - "King" and "Queen" are as unrelated as "King" and "Table."
 - Order and context of words are lost.
- **Meaning captured?** No real semantic understanding—only presence/absence or frequency.

Word2Vec

- **How it works:** Learns **dense vector representations** (embeddings) of words using neural networks.
- **Key idea:** Words that appear in **similar contexts** have similar meanings.
- **Two main models:**

Name :- Prathamesh Arvind Jadhav

Roll No:- 4059

- **CBOW (Continuous Bag-of-Words):** Predicts a word from surrounding context.
- **Skip-Gram:** Predicts surrounding context from a given word.
- **Why it captures meaning:**
 - By training on large text corpora, Word2Vec places semantically similar words closer in vector space.

Q4. What are the advantages of Word2Vec over TF-IDF?

Ans->

- **Advantages of Word2Vec over TF-IDF**

1. **Captures Semantic Meaning**

- Word2Vec places semantically similar words near each other in vector space.
- TF-IDF treats each word as independent, so synonyms aren't recognized.

2. **Dense Representation**

- Word2Vec produces low-dimensional vectors (e.g., 100–300 dimensions), making computations faster and more memory-efficient.
- TF-IDF creates very high-dimensional sparse vectors (one dimension per vocabulary word).

3. **Handles Synonyms & Context**

- Word2Vec recognizes that *"good"* and *"excellent"* are similar in meaning.
- TF-IDF cannot — both get totally separate vector positions.

4. **Word Relationships & Analogies**

- Word2Vec captures relationships like *"Paris - France + Italy = Rome"*.
- TF-IDF has no way of encoding such relationships.

5. **Better for Downstream Tasks**

- Word2Vec embeddings improve performance in tasks like sentiment analysis, machine translation, or recommendation systems.
- TF-IDF is mostly useful for keyword-based tasks like search or text classification.

Name :- Prathamesh Arvind Jadhav

Roll No:- 4059

6. **Generalization Beyond Exact Words**

- Word2Vec can relate unseen but contextually similar words if trained on enough data.

• **Conclusion:-**

This experiment demonstrates how text can be represented numerically using Bag-of-Words, TF-IDF, and Word2Vec embeddings. While BoW and TF-IDF provide sparse representations, Word2Vec generates dense semantic embeddings that capture contextual meaning.